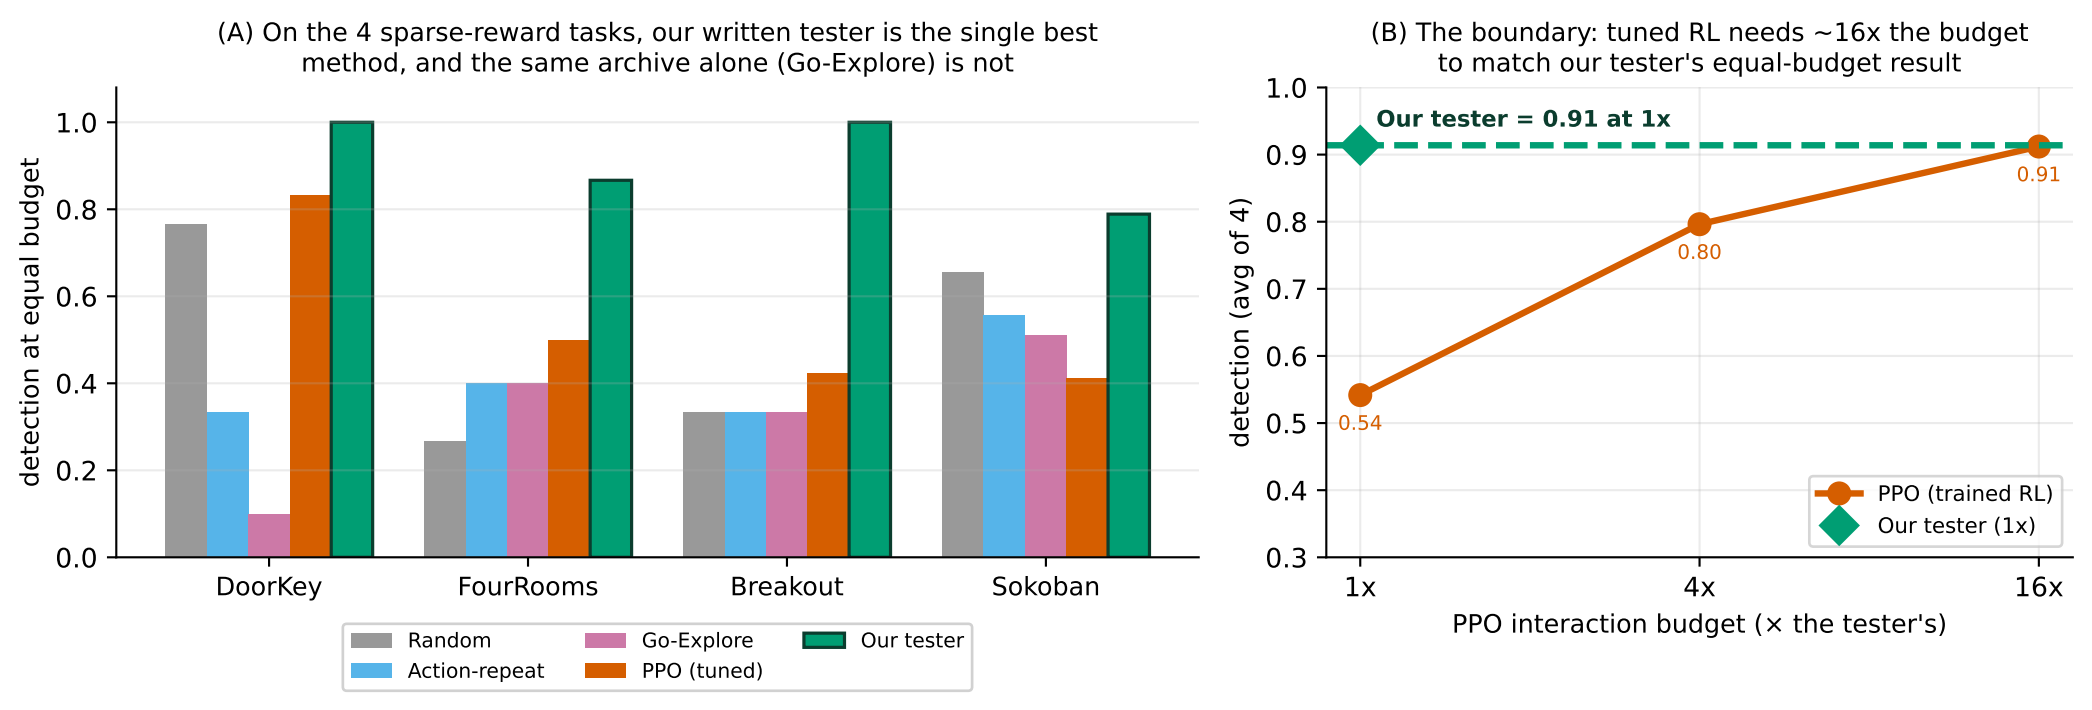
(A) On the 4 sparse-reward tasks, our written tester is the single best method, and the same archive alone (Go-Explore) is not
detection at equal budget
0.0
0.2
0.4
0.6
0.8
1.0
DoorKey
FourRooms
Breakout
Sokoban
Random
Action-repeat
Go-Explore
PPO (tuned)
Our tester
(B) The boundary: tuned RL needs ~16x the budget to match our tester's equal-budget result
detection (avg of 4)
0.3
0.4
0.5
0.6
0.7
0.8
0.9
1.0
Our tester = 0.91 at 1x
0.54
0.80
0.91
PPO (trained RL)
Our tester (1x)
1x
4x
16x
PPO interaction budget (× the tester's)