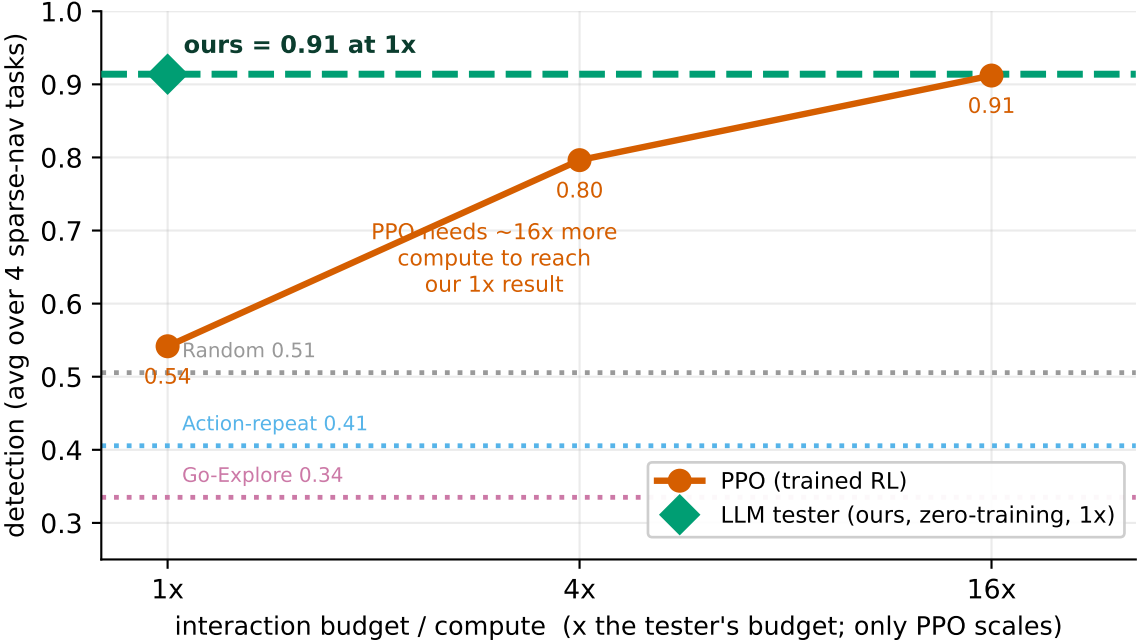
At equal budget the zero-training tester dominates; trained RL needs ~16x the compute to catch up, baselines never do
ours = 0.91 at 1x
0.91
0.80
PPO needs ~16x more compute to reach our 1x result
0.54
Random 0.51
Action-repeat 0.41
Go-Explore 0.34
PPO (trained RL)
LLM tester (ours, zero-training, 1x)
detection (avg over 4 sparse-nav tasks)
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.3
1x
4x
16x
interaction budget / compute (x the tester's budget; only PPO scales)